


Systematic Failure of Activation Steering on Llama-3.1-8B-Instruct

Each approach was tested, failed, and informed the next experiment



CONCLUSION

Llama-3.1-8B-Instruct is resistant to activation-based steering.
Even patching ALL attention heads with neutral activations increases sycophancy.

KEY INSIGHT

RL-optimized vectors (GRPO) work on this model. Mean-difference methods (CAA, CCM) do not.
The difference is optimization: extracted vectors capture correlation, not causation.